The EM Algorithm

Ritwik Sahani/ Aayush Goyal

IIT Hyderabad

February 28, 2019

Suppose we have some data points $x_1 \ldots\ldots x_n$

and we also tell you that these points belong to two different probabilistic distributions, say, two gaussians such that we know which point belogs to which particular distribution.

Can we find the parameters μ_b, σ^2 for the two gaussians?

Indeed, we can.

$\mu = \sum_{i=1}^{n_b} \frac{x_i}{n_b}$
$\sigma^2 = \sum_{i=1}^{n_b} \frac{(x_i - \mu_b)^2}{n_b}$
and similarly for the other distribution.

Now, suppose data points are there but we don't know the as to which gaussian a particular point belongs.

But what we do know is that these come from two gaussians, whose parameters μ, σ^2 is known

Can we, for each point, decide which of the two gaussians it is more likely to belong to?
Yes, this can be done by claculating a posteriori prob.

$$P(b|x_i) = \frac{P(x_i|b)P(b)}{P(x_i|b)P(b)+P(x_i|a)P(a)}$$

where $P(x_i|b)$ can be obtained from the gaussian $N(\mu_b, \sigma_b^2)$

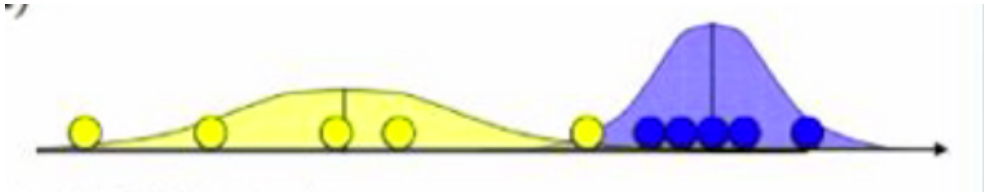

But what happens, if we neither knew the source gaussian, nor their parameters?

Then we use the EM algorithm method, an iterative method, to find the same.

Step-1: Expectation

Creates a function for the expectation of the log-likelihood evaluated using the current estimate for the parameters.

⋆ **Basically we will assume some initial parameters and assign each point with its a- posteriori probability**

Step-2: Maximization

Which computes parameters maximizing the expected log-likelihood found on the E step. These parameter-estimates are then used to determine the distribution of the latent variables in the next E step.

⋆ **Basically, we will again calculate the parameters for the next E -step**

Two-Component Mixture Model

We consider a simple mixture model for density estimation, and the corresponding EM algorithm for carrying out maximum likelihood estimation.

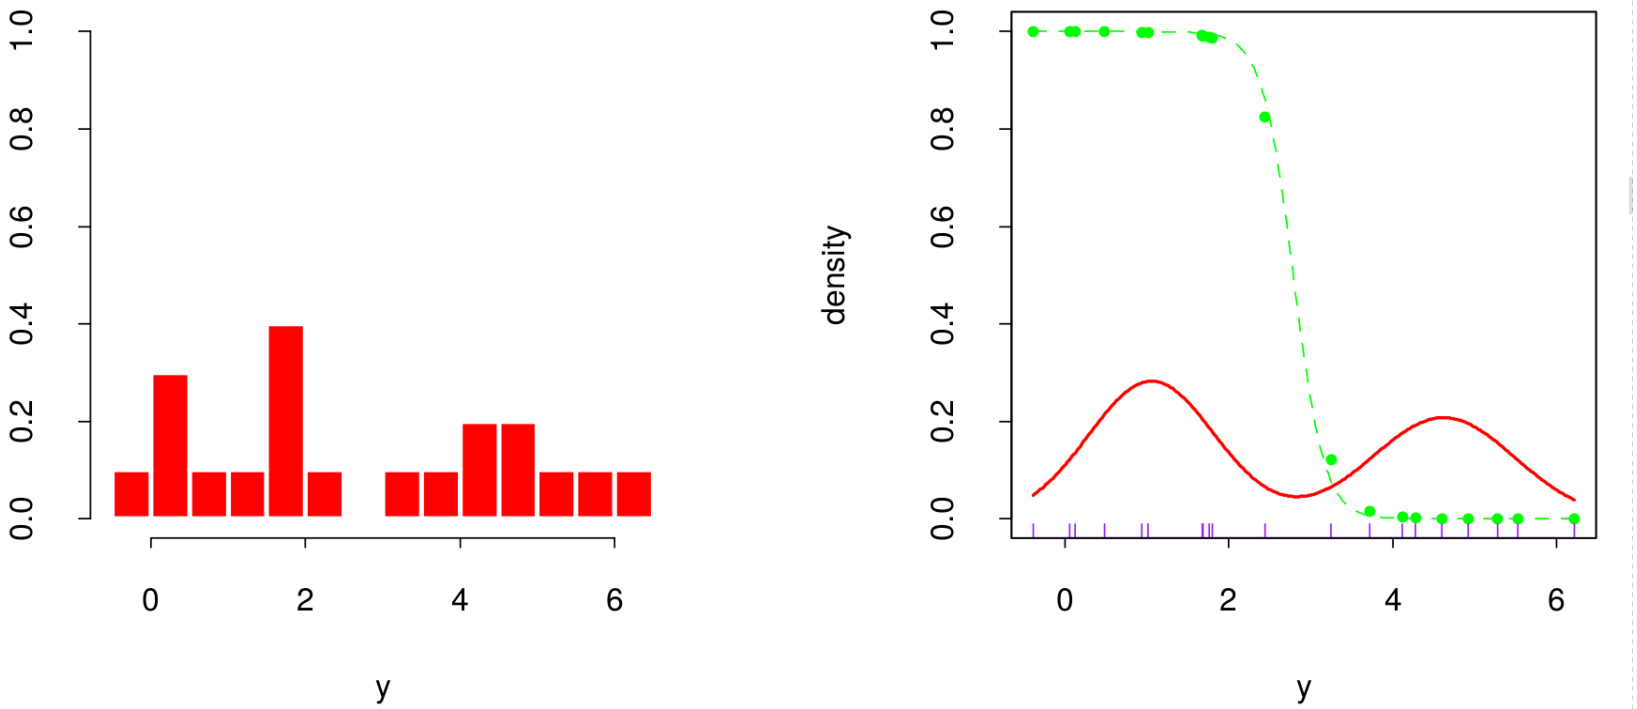

Figure: Mixture example. (Left panel:) Histogram of data. (Right panel:) Maximum likelihood fit of Gaussian densities (solid red) and responsibility (dotted green) of the left component density for observation y, as a function of y.

TABLE 8.1. *Twenty fictitious data points used in the two-component mixture example in Figure 8.5.*

-0.39	0.12	0.94	1.67	1.76	2.44	3.72	4.28	4.92	5.53
0.06	0.48	1.01	1.68	1.80	3.25	4.12	4.60	5.28	6.22

Due to Bimodalty of data Gaussian distribution would be inappropriate. So we will model Y as a mixture of two distinct distributions.

$Y_1 \equiv N(\mu_1, \sigma_1^2)$
$Y_2 \equiv N(\mu_2, \sigma_2^2)$
$Y = (1 - \Delta) * Y_1 + \Delta * Y_2$
where $\Delta \in \{0, 1\}$ with $Pr(\Delta = 1) = \pi$

This generative representation is explicit: generate a $\Delta \in \{0, 1\}$ with probability π, and then depending on the outcome, deliver either $Y_1 or Y_2. Let \phi\ \theta$ (x) denote the normal density with parameters $\theta = (\mu, \sigma_2)$. Then the density of Y is $gY(y) = (1 - \pi)\phi\theta_1(y) + \pi\phi\theta_2(y)$.

Now suppose we wish to fit this model to the data in Figure 8.5 by maximum likelihood.The parameters are
$\theta = (\pi, \theta_1, \theta_2) = (\pi, \mu_1, \sigma_{12}, \mu_2, \sigma_{22})$
The log-likelihood based on the N training cases is

$l(\theta; Z) = \sum_{i=1}^{N} log[(1-\pi)\phi\theta_1(y_i) +$
pi$\phi\theta_2(y_i)]$

Direct maximization of $l(\theta; Z)$ is quite difficult numerically, because of the sum of terms inside the logarithm. There is, however, a simpler approach. We consider unobserved latent variables θ_i taking values 0 or 1 as in (8.36): if $\Delta_i = 1$then Y_1 comes from model 2, otherwise it comes from model 1.

$l_0(\theta; Z, \Delta) = \sum_{i=1}^{N}[(1-\Delta_i)log\phi_{\theta_1}(y_i) + \Delta_i log\phi_{\theta_2}(y_i)] +$
$\sum_{i=1}^{N}[(1-\Delta_i)log(1-\pi)(y_i) + \Delta_i log\pi]$

Now we apply the E-M Algorithm for Two-component Gaussian mixture

1. Take initial guesses for the parameters $\hat{\mu}_1, \hat{\sigma}_1^2, \hat{\mu}_2, \hat{\sigma}_2^2, \hat{\pi}$
2. *Expectation Step*: compute the responsibilities

$$\hat{\gamma}_i = \frac{\hat{\pi}\phi_{\hat{\theta}_2}(y_i)}{(1-\hat{\pi})\phi_{\hat{\theta}_1}(y_i) + \hat{\pi}\phi_{\hat{\theta}_2}(y_i)}, \ i = 1, 2, \ldots, N.$$

3. *Maximization Step*: compute the weighted means and variances:

$$\hat{\mu}_1 = \frac{\sum_{i=1}^N (1-\hat{\gamma}_i) y_i}{\sum_{i=1}^N (1-\hat{\gamma}_i)}, \qquad \hat{\sigma}_1^2 = \frac{\sum_{i=1}^N (1-\hat{\gamma}_i)(y_i - \hat{\mu}_1)^2}{\sum_{i=1}^N (1-\hat{\gamma}_i)},$$

$$\hat{\mu}_2 = \frac{\sum_{i=1}^N \hat{\gamma}_i y_i}{\sum_{i=1}^N \hat{\gamma}_i}, \qquad \hat{\sigma}_2^2 = \frac{\sum_{i=1}^N \hat{\gamma}_i (y_i - \hat{\mu}_2)^2}{\sum_{i=1}^N \hat{\gamma}_i},$$

   and the mixing probability $\hat{\pi} = \sum_{i=1}^N \hat{\gamma}_i / N$.

4. Iterate steps 2 and 3 until convergence.

Note that the actual maximization of the likelihood occurs when we put a spike of infinite height at any one data point, that is, $\hat{\mu}_1 = y_i$ for some i and $\hat{\sigma}_1^2 = 0$. This gives infinite likelihood, but is not a useful solution. Hence we are actually looking for a good local maximum of the likelihood, one for which $\hat{\sigma}_1^2, \hat{\sigma}_2^2 > 0$. To further complicate matters, there can be more than one local maximum having $\hat{\sigma}_1^2, \hat{\sigma}_2^2 > 0$.. In our example, we ran the EM algorithm with a number of different initial guesses for the parameters, all having $\hat{\sigma}_k^2 > 0.5$, and chose the run that gave us the highest maximized likelihood

Iteration	$\hat{\pi}$
1	0.485
5	0.493
10	0.523
15	0.544
20	0.546

Figure: Selected iterations for E-M algorithm

The final maximum likelihood estimates are $\hat{\mu}_1 = 4.62$ $\hat{\sigma}_1^2 = 0.87$, $\hat{\mu}_2 = 1.06$,$\hat{\sigma}_2^2 = 0.77$, $\hat{\pi} = 0.546$.

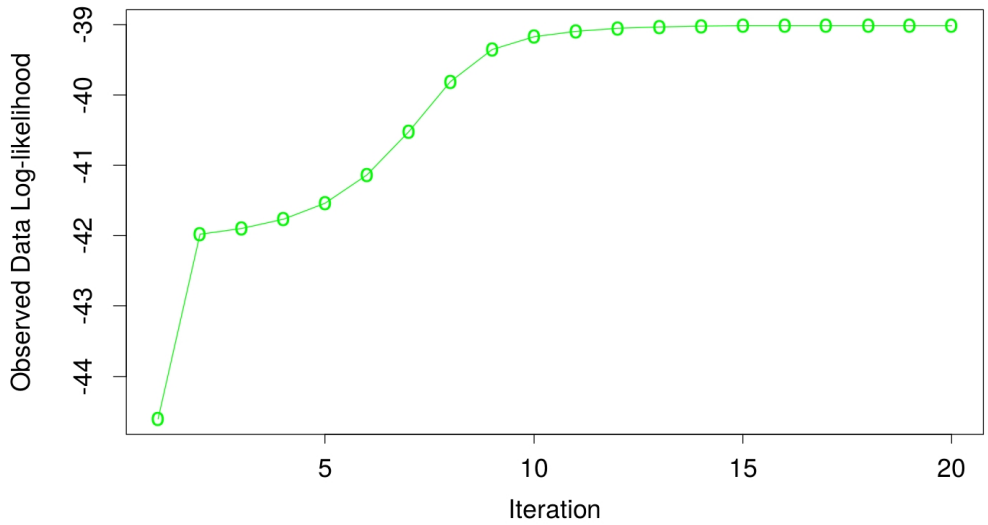

Figure: observed data log-likelihood as a function of the iteration number

EM Algorithm in general

1. Start with initial guesses for the parameters $\hat{\theta}^{(0)}$.
2. *Expectation Step*: at the jth step, compute

$$Q(\theta', \hat{\theta}^{(j)}) = \mathrm{E}(\ell_0(\theta'; \mathbf{T})|\mathbf{Z}, \hat{\theta}^{(j)})$$

   as a function of the dummy argument θ'.
3. *Maximization Step*: determine the new estimate $\hat{\theta}^{(j+1)}$ as the maximizer of $Q(\theta', \hat{\theta}^{(j)})$ over θ'.
4. Iterate steps 2 and 3 until convergence.

References

1) The Elements of Statistical Learning - Trevor Hastie, Robert Tibshirani, Jerome Friedman

2) https://en.wikipedia.org/

3) Victor Lavrenko - Youtube